

Report: Assignment 1 - Basic

Tanjina Islam (2609513, vu-netID:tim500), Konrad Karas (2624767, vu-netID:kks206) and Izak de Kom (2629935, vu-netID:ikm700), group 27

22 April 2018

Task 1: Explore a small dataset

Data preprocessing

We kept the following features in our preprocessed dataset:

```
data = read.csv('ODI-2018_clean.csv')
colnames(data)

## [1] "What.programme.are.you.in."
## [2] "Have.you.taken.a.course.on.machine.learning."
## [3] "Have.you.taken.a.course.on.information.retrieval."
## [4] "Have.you.taken.a.course.on.statistics."
## [5] "Have.you.taken.a.course.on.databases."
## [6] "What.is.your.gender."
## [7] "Chocolate.makes.you....."
## [8] "birth_day"
## [9] "birth_month"
## [10] "birth_year"
## [11] "Give.a.random.number"
## [12] "bedtime"
## [13] "good_day_1"
## [14] "good_day_2"
```

The study program feature was cleaned using the 'format_study_program.R' script. The course features were relatively clean already and did not need additional cleaning. This was also the case for the 'gender', 'chocolate makes you' and 'give a random number' features. The original 'birthday' feature was cleaned using the 'clean_birthday_bedtime.py' script and splitted into a day, month and year feature. Unfortunately, the 'bedtime' feature cleaning was problematic and was manually formatted. Finally, the 'good day' features were cleaned with the 'GoodDay_cleanup.R' script using exact string matching and the levenstein distance coefficient.

Exploratory data analysis

The cleaned data consisted of 218 samples and 14 features:

```
nrow(data)

## [1] 218
```

```
ncol(data)
```

```
## [1] 14
```

More men participate in the data mining techniques course:

```
pie(table(data$What.is.your.gender.), col = rainbow(3), labels = c("Unknown: 4", "Females: 63",
```




Figure 1. Pie chart showing the gender distribution in the data mining techniques course.

The birth years distribution (left out two students who were born in 1768 and 1931):

```
plot(table(data$birth_year), xlab = "Year", ylab = "Number of students", xlim = c(1981,2000))
```




Figure 2. Birth year distribution in the data mining techniques course.

A short description of the answers of the ‘Chocolate makes you...’ question. First, the data is converted in a frequency table using the `table` function. Then the `names` are redefined to define

shorter class names. Finally, the data is plotted in figure 2.

```
choc = table(data$Chocolate.makes.you.....)
names(choc) = c("NA", "fat", "No idea", "Neither", "slim", "unknown")
barplot(choc)
```




Figure 3. Chocolate makes you... Answer frequencies.

Basic classification/regression

A simple reression using gender as dependent variable and the ‘chocolate makes you...’ as independent variable:

```
library(boot)
glm = glm(data$What.is.your.gender.~data$Chocolate.makes.you....., family = 'binomial')
summary(glm)
```

```
##
## Call:
## glm(formula = data$What.is.your.gender. ~ data$Chocolate.makes.you.....,
##      family = "binomial")
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.409e-06  2.409e-06  2.409e-06  2.409e-06  2.409e-06
##
## Coefficients:
##
##                                Estimate
## (Intercept)                   -26.57
## data$Chocolate.makes.you.....fat    53.13
## data$Chocolate.makes.you.....I have no idea what you are talking about  53.13
```

```

## data$Chocolate.makes.you.....neither      53.13
## data$Chocolate.makes.you.....slim          53.13
## data$Chocolate.makes.you.....unknown       53.13
##                                             Std. Error
## (Intercept)                               356124.00
## data$Chocolate.makes.you.....fat           358522.17
## data$Chocolate.makes.you.....I have no idea what you are talking about 359234.31
## data$Chocolate.makes.you.....neither       358658.72
## data$Chocolate.makes.you.....slim          368623.31
## data$Chocolate.makes.you.....unknown       436160.74
##                                             z value
## (Intercept)                               0
## data$Chocolate.makes.you.....fat           0
## data$Chocolate.makes.you.....I have no idea what you are talking about 0
## data$Chocolate.makes.you.....neither       0
## data$Chocolate.makes.you.....slim          0
## data$Chocolate.makes.you.....unknown       0
##                                             Pr(>|z|)
## (Intercept)                               1
## data$Chocolate.makes.you.....fat           1
## data$Chocolate.makes.you.....I have no idea what you are talking about 1
## data$Chocolate.makes.you.....neither       1
## data$Chocolate.makes.you.....slim          1
## data$Chocolate.makes.you.....unknown       1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 1.2764e+01 on 217 degrees of freedom
## Residual deviance: 1.2647e-09 on 212 degrees of freedom
## AIC: 12
##
## Number of Fisher Scoring iterations: 25
# cross validation:
cv.glm(data[,c(6,7)],glm)

## $call
## cv.glm(data = data[, c(6, 7)], glmfit = glm)
##
## $K
## [1] 218
##
## $delta
## [1] 0.009132228 0.009132228
##
## $seed
## [1] 403 1 766510385 508638563 812504882
## [6] 1378222212 -433234745 2063814557 1769831188 1066600890

```

##	[11]	-1030447099	-1803758113	-1237407642	-651520016	-573578989
##	[16]	989937329	-1688405664	2140622590	68496009	-1280293445
##	[21]	-1054556278	-1032935828	102045679	1213313061	924664028
##	[26]	2049124018	-568833619	632765799	-165451122	-483167384
##	[31]	1675406187	504796425	-1175688296	1443490502	923455521
##	[36]	354194611	-1622047198	1776302100	-1295795049	-654633331
##	[41]	689600420	-899782262	1857770997	1153878095	688119638
##	[46]	1013822976	-250804413	331013857	1107377232	1624605870
##	[51]	1829347417	-394346389	-2029226534	131653276	-837488481
##	[56]	-1646490891	852301708	1791637090	376146493	1234154679
##	[61]	1493035870	904973976	1137356283	-1252724839	-2103556056
##	[66]	148484310	1749448017	-37805053	1529957266	1031252260
##	[71]	417492967	1675949117	1748222644	-824815334	-766827931
##	[76]	-16185601	-38634618	1511841040	1160107443	-1884453295
##	[81]	753788416	-1151524130	-125718231	1088330971	-1018024982
##	[86]	947092364	1737374927	-767957883	1846849340	5451282
##	[91]	1910910221	-1149864377	661728302	-452906616	642771083
##	[96]	1820295785	721419832	-1657825562	-1892502911	1222852115
##	[101]	1017268674	-1767087500	750879863	802704749	1873932804
##	[106]	-197275350	469173909	-887984849	192010294	-1933775776
##	[111]	-740326749	-1474841151	-1013808720	1579999438	694156857
##	[116]	-83738549	1184297466	-896193284	1318371647	-1729315307
##	[121]	722304300	1920093698	1974260957	1137430487	-319444290
##	[126]	-1603162888	376115931	1877985593	279743624	-1998717898
##	[131]	530434161	-565282141	231268338	430961476	-590586873
##	[136]	-427951651	799229652	-1195218438	2053621061	-1028845153
##	[141]	14955302	1723927472	1625487443	-453872783	1861616288
##	[146]	21689918	744815945	53691387	-697161782	231151916
##	[151]	1832236335	-57132059	852132508	-310917390	-1032857747
##	[156]	-601382745	-2080103730	-1957081048	1834576939	-1796710839
##	[161]	-2053007528	-1258028538	152710113	1988544883	888072674
##	[166]	-236679852	-864279465	-131659699	2329572	-849024566
##	[171]	479424565	-129700209	727433110	-561790144	1034415619
##	[176]	-1228803039	-960039152	-759337746	-698076007	-1771462357
##	[181]	736385050	-1513336100	-58125601	1866878133	432122444
##	[186]	-294103390	-589382147	-1347135753	-623624162	-618704168
##	[191]	-1602619845	-1873919015	1863987560	-1264957290	-883620335
##	[196]	-571555901	-448756782	-1853194652	1012764071	1557734141
##	[201]	1094096628	-1080016422	-2075065051	773653567	-983485498
##	[206]	-1654251440	875170931	-1062657903	277971136	-284595554
##	[211]	478324329	754882203	-81462358	904749004	647402127
##	[216]	8633029	-2091724164	695664338	1811536461	-1700710649
##	[221]	1555684078	-725217336	863310027	186282281	43403640
##	[226]	804658854	336840897	1793409363	-1217647102	-1680911044
##	[231]	1684868210	-2091521936	502104100	1127992376	820162330
##	[236]	-323261232	1216397116	168814484	-1708764286	849037760
##	[241]	-78873028	886293904	1533099554	1328456968	1237003916
##	[246]	-537092324	-210714478	2117777792	-1675598316	-1705795800

## [251]	2026183226	-203735600	381757804	1017250740	979523602
## [256]	374732256	-2128951348	-1797136160	1682216994	1265525032
## [261]	1979162124	-1652971716	505070578	2073697008	1668566596
## [266]	318751320	-1825878214	-1810897424	-2097195332	-1355237100
## [271]	1396720322	-2041023392	680264572	208934864	131555618
## [276]	-177472856	2077596556	-875644740	-445518030	2012064320
## [281]	503512244	1049436744	-1867323974	-1482653168	535622380
## [286]	-1007181484	1930477842	-1044722272	361538252	-1682629024
## [291]	-1880590398	780887848	560951212	-519380164	-561796750
## [296]	-635110608	1976233252	-1578521800	541892314	-687265904
## [301]	-1266790660	1768270740	-38231998	1892537600	341541948
## [306]	-1141325488	-1408921438	-665657016	-749398260	246006492
## [311]	714173778	948683904	457711700	-1467215896	-1590810374
## [316]	-1662244016	384433964	-375838348	-1898386542	-959141792
## [321]	-940651124	-625321120	730216930	-1090817112	-770296628
## [326]	381091836	529130930	-319752720	-904033340	-657513128
## [331]	24009082	-1940716624	-1603696196	-1436600236	-1225410622
## [336]	-506712032	-96520004	1693160016	-1832078366	1548517480
## [341]	266093644	-892921604	-1425241102	-1200305920	-313835148
## [346]	-1314264056	-147816518	-1186057776	297303020	-818142828
## [351]	1103434194	1892553056	-335425716	-2120197856	1849343874
## [356]	-870928664	-1125170964	-1685638084	-598525070	-1151975824
## [361]	1467061284	528239416	-2009551718	-1635052848	-904691524
## [366]	-1487327596	-6801790	-1002092352	1296180028	-800531312
## [371]	-937361886	-723136504	573545740	-592040804	954531346
## [376]	287997824	1189629460	-944526808	-997100230	-1936307760
## [381]	-369675156	22358324	-1964271982	191990624	-21957812
## [386]	-661243936	-1158441054	1137018792	-79024756	-1937602372
## [391]	1968769394	-619997968	-1404424124	1961075032	-1316673990
## [396]	-1197900944	-1003997124	-289906924	-2099157566	-1683784992
## [401]	-1615685764	-1364402480	-996452190	103779368	-1347427060
## [406]	1842767804	-413177550	-875796032	-2091496908	806362056
## [411]	1252997306	-366171760	-1720495636	794280532	-548999022
## [416]	1761324832	-178722356	1450546656	-1462307134	-24845272
## [421]	-575819092	1100997820	-1630475790	624210224	-335244124
## [426]	-582449480	1925299546	-784042608	-548318340	1828579604
## [431]	1482163010	1649203968	-602114116	-797319344	-1432992990
## [436]	-115548856	-89298420	-1385707428	357137618	1532566016
## [441]	2087185748	-1641768472	2034928634	1954660560	1726123308
## [446]	-986471692	-919660142	-897797024	-290413300	-401927456
## [451]	2013415778	-971695448	157942348	-1199672964	102778162
## [456]	825946224	-1020943028	-273840867	-308199257	-1414516784
## [461]	1540450158	-534868453	-1700464771	-596869782	642110472
## [466]	747166897	1809316643	2015823908	1293290194	-599335129
## [471]	1808367985	-1093089354	-426656940	33408485	667828207
## [476]	1060806440	-1575311482	-486763597	-193858795	-792392334
## [481]	1774219712	-244096471	1127458459	817590124	-1429768326
## [486]	-277344081	965847993	-229589554	-1923651300	-1560224627

```

## [491] 2115522807 -345875456 1956273758 242467115 -1840340435
## [496] 1094591130 -659268008 -1062673983 -1407783309 -1990988908
## [501] 1079501346 -1662718665 -46429183 -675041018 -94120636
## [506] 2014836853 1700452831 -965745352 265844502 1706260259
## [511] -67894043 -2019835646 -335498384 -319739111 -51000565
## [516] -1362225156 -1506380758 235608159 1626460137 -2081672962
## [521] -1731436308 468845629 -886474745 -160398352 -177975410
## [526] -1212723781 -272684131 419934538 1716246888 1404358481
## [531] 517786691 -697330364 -1097681422 130072519 -416480879
## [536] 1928770838 756960628 -2006430715 1859533903 -1510911928
## [541] -1201093146 396368659 -92570507 -1574705582 -1732665184
## [546] 1345485193 -1471481413 1848432140 -756756582 -1530300657
## [551] -2093290151 588318702 883486332 -319367443 -99442281
## [556] 1996024416 -599744450 -293803765 -1445270643 18229434
## [561] -12194056 904766625 -89391021 500706164 1353636098
## [566] 543586711 1884152673 -1527236698 -63925852 764299349
## [571] -2059346689 1234264472 -1299547466 115184451 -1204744571
## [576] -524106974 -84477936 -516799559 1676551275 495875932
## [581] -641446262 -1036608897 -915012663 1102159966 -1876540660
## [586] 1070754525 -1189444633 -1261283312 854209838 892340699
## [591] -1270624579 23870506 -2040054584 829413873 -884828317
## [596] -620310684 -1389844846 1465746279 -511365839 -744978826
## [601] -1804634220 -614923355 347059375 -582686360 -846327482
## [606] 1291348211 -2104592811 1960188466 1370957824 1407040745
## [611] 1909537627 1818451372 -1515600454 1579582319 740709625
## [616] 1547618702 -1572979620 -612528435 -7132617 998596800
## [621] 674066334 141309035 -723020819 1630385754 -1354991336
## [626] 1295112069

```

There seems to be no clear relationship between the answer type and the gender.

Task 2: Kaggle Competition: Titanic: Machine Learning from Disaster

The sinking of the RMS Titanic is one of the most infamous shipwrecks in history. On April 15, 1912, during her maiden voyage, the Titanic sank after colliding with an iceberg, killing 1502 out of 2224 passengers and crew. This sensational tragedy shocked the international community and led to better safety regulations for ships.

The aim of the competition is to predict who among the passengers and crew was more likely to survive than others. Kaggle provides two datasets: *train* and *test*. While both datasets reference to passengers details, only *train* dataset contains information if passenger survived or not. Our goal is to predict which passenger from *test* dataset survived the sinking of Titanic.

Preparation

As mentioned before, Kaggle has provided two separate datasets: *train* and *test*. Both datasets contain details about passengers and their trip. The only difference between them is *Survived* column in *train* dataset that indicates if passenger has survived the Disaster. That dataset will be used to train our models.

Before attempting to perform predictions, we focused on given data and tried to retrieve more interesting facts based on Feature Engineering. Because the only column that differ is *Survival*, for further processing we decided to datane both datasets into big one. Such an approach allows us to perform more adequate data analyse as we have a full insight of traveling passengers.

Data exploration

To have a better insight into assignment, we had to explore given data. That's the most important step - we have to be aware of all the details to work efficiently. Whole dataned dataset contains 1309 records (passengers) with 12 variables. In this part we will take a closer look to every attribute.

In datasets we can distinguish several (12) columns:

- Survived - indicated if given passenged survived
- PassengerId - passenger index in dataset
- Pclass - the ticket class (1,2,3)
- Name - full name of passenger, including their title
- Sex - sex of passenger
- Age - age of passenger
- SibSp - number of siblings or spouses traveling with passenger
- Parch - number of parents or childern traveling with passenger
- Ticket - ticket number
- Fare - passenger fare
- Cabin - passenger's cabin number
- Embarked - port of embarkation (C = Cherbourg, Q - Queenstown, S = Southampton)

Lets make a closer look into several variables.


Name

In given dataset we can see that *Name* attribute contains string with passenger's name, surname and title.

example: *Allison, Master. Hudson Trevor*


Fortunately, all rows in *Name* column follow the same string pattern (*surname, title first name*). Thanks to this fact, we will be able to retrieve more additional infrmation about passengers, like common surnames or titles.

Sex



Investigating Sex attribute we can see that there were 466 females and 843 males onboard. That gives us the first easy grouping of passengers.

Age



Regarding Age attribute, we can see that this variable varies up to 80 with mean around 23

Feature Engineering

After investigation of given dataset we can distinguish columns that seem to be useful for further processing to retrieve even more data. In such an approach we are able to create additional columns

with relevant variables that could result in better prediction accuracy.

Feature: Title

As mentioned before, Name column contains not only name and surname of passenger but also a title (like Sir., Mr., Mrs., ...). Following common pattern (*surname, title first name*) we can retrieve additional Column in our dataset that would group our passenger by Title. In addition, groups of unique similar titles were replaced by the same variable (like 'Capt', 'Don', 'Major', 'Sir' => 'Sir').

As a result we obtained a fixed set of values with 11 levels.

Feature: Family

Basing on variables *SibSp* (number of siblings or spouses), *Parch*(number of parents or child) and Surnames retrieved from *Name* variable we are able to group passengers by families. Assuming that during disaster, every person takes care about their relatives, we think that it can be a significant factor in predictions.

Our assumptions:

- the number of relatives with who each passenger was traveling is calculated as follows: $SibSp + Parch + 1$ - result is family size
- if family size is less or equal 2 we assume that the value is not relevant and we mark such a family as n/a

As a result we obtained Family attribute with 97 levels.

Feature: Deck

Analysing *Cabin* attribute we figured out that each cabin number consists of Deck Level and Room number (like C40 => Deck C, Room 40). Because Deck Level could play important role in evacuation, we assumed that it's a significant attribute. We decided to create a new attribute called Deck and we assigned relevant Deck Level to each passenger. Unfortunately, not every passenger had a Cabin number assigned, in such a case we marked Deck as 'U'.

Result:

##	A	B	C	D	E	F	G	T	U
##	22	65	94	46	41	21	5	1	1014

Feature: TicketType

Looking into ticket numbers we can see that some tickets have common prefix that could refer to Ticket Type of place of purchase (example: STON/02 42342). We decided to retrieve that ticket prefix and create a new attribute for each passenger. If ticket didn't have any prefix, we marked TicketType as 'num'.

As a result we obtained TicketType factor with 51 levels.

Missing values

We have found that some records lack in Age attribute. In such a situation we decided to use a Decision tree to predict missing Age values. As significant factors we marked attributes: Pclass, Sex, FamilySize, Embarked, Title, SibSp, Parch.

Also Fare column had some missing values. In such a case we replaced missing values with median of all ticket Fares.

Classification and evaluation

By analysing our data and engineering some additional features we have enriched our dataset.

Within all columns we decided that only few of them play significant role in predictions.

Chosen factors: *Pclass, TicketType, Sex, Deck, Age, SibSp, Parch, Fare, Embarked, Title, FamilySize, FamilyID*

Creating a setup

To evaluate classifiers we will need to create a proper setup. In this case we decided to use *train* data from Kaggle as it contains *Survived* column. For evaluation purposes we decided to split the data for training and testing sets with ratio 70/30 (70% - training, 30% - testing). While splitting the data we based on random ordering.

For evaluation we decided to use two non-linear algorithms: k-Nearest Neighbour Classification and Conditional inference trees. Both classifiers were trained and tested with the same sets of data. For evaluation analysis we used Confusion Matrix.

Factors that we took into account:

- Accuracy - how well results were predicted
- 95 CI - confidence intervals, our final score should match into calculated intervals
- Kappa - accuracy through random predictions
- F1 - model that takes recall and precision into account

Evaluation of k-Nearest Neighbour Classification

Accuracy : 0.6929

95% CI : (0.6338, 0.7477)

Kappa : 0.3338

F1 : 0.5638

Evaluation of Conditional inference trees

Accuracy : 0.809

95% CI : (0.7566, 0.8543)

Kappa : 0.5929
F1 : 0.7437

Kaggle Submission

For Kaggle competition we decided to use Conditional inference trees as it gives us higher results in included evaluation factors.

We have submitted our Prediction in Kaggle system and obtained satisfactory result 0.82296 which is top 3% in leaderboard (username VUDM27). This result also matches into expected Confidence Intervals calculated during evaluation.

Task 3: Research and theory

Task 3.A – Research: State of the art solutions (10 points)

Data mining competition: “Planet: Understanding the Amazon from Space”, July 2017

In this data mining competition, the participants were provided with a dataset consisting of over 40,000 satellite images. Each image had a resolution of 256 by 256 pixels, and covered an area of approximately 950 by 950 meters. The satellite images were captured above the Amazon in South-America. The goal of the competition was to successfully classify these satellite images. Each image could be classified into multiple classes:

- **Atmospheric conditions:** clear, partly cloudy, cloudy, and haze
- **Common land cover and land use types:** rainforest, agriculture, rivers, towns/cities, roads, cultivation, and bare ground
- **Rare land cover and land use types:** slash and burn, selective logging, blooming, conventional mining, artisanal mining, and blow down

The participants’ model predictions were evaluated with their mean F2 score. This score determines the model prediction accuracy using the precision and recall measures. The formula of the score:

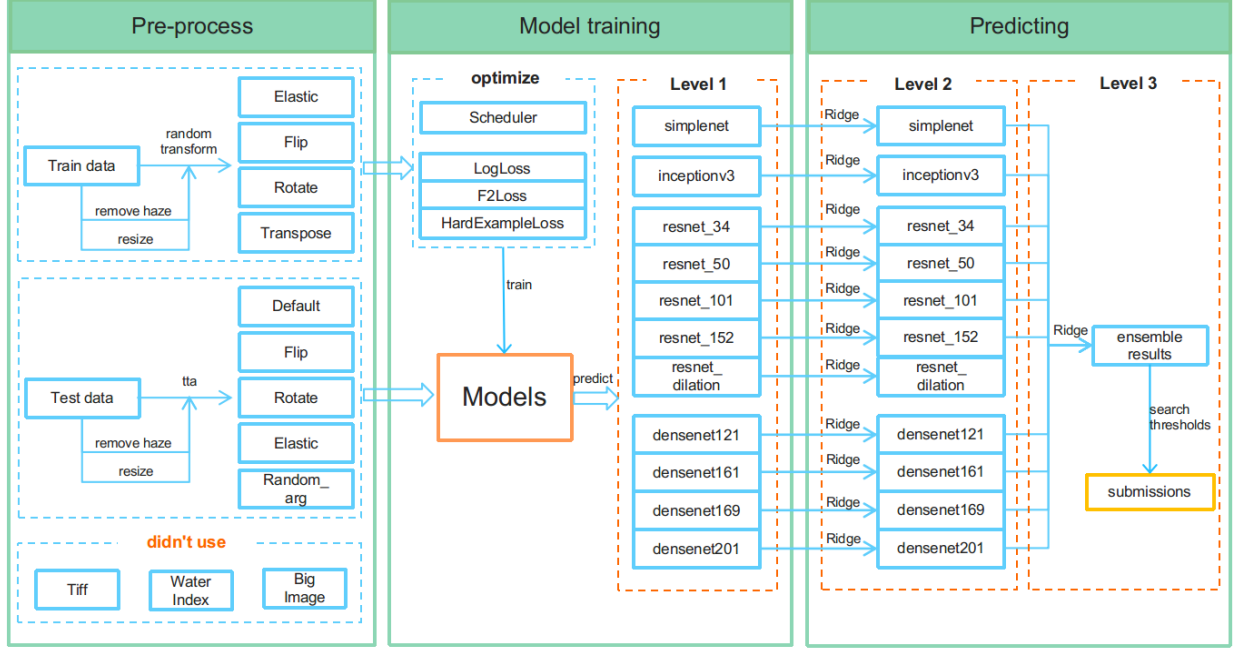
$$(1 + \beta^2) \frac{pr}{\beta^2 p + r} \text{ where } p = \frac{tp}{tp + fp}, \quad r = \frac{tp}{tp + fn}, \quad \beta = 2.$$

The winner of the competition was a Kaggle user named “bestfitting”. He managed to win the competition by fine-tuning 11 convolutional neural networks (CNN’s), and using these CNN’s to build an ensemble model. This ensemble model was used to predict the final classes.

The pre-processing step consisted of resizing the satellite images, removing haze from the images,

and augmenting the data (e.g. flipping, rotating and transposing). Next, one *simplenet*, one *inceptionv3*, five *resnet* and four *densenet* CNN's were trained on the labeled training data. The resulting models were then ensembled by using a ridge regression model, allowing for the selection of the strongest models for each label prediction.

The main reasons why this model won the competition was the creation of an ensemble model. As different models had different capabilities on each class label, the combination of the models resulted in a higher accuracy.



Task 3.B – Theory: MSE verse MAE

Mean Squared Error(MSE)

The mean squared error (MSE) of an estimator measures the average of the squares of the errors that is, the difference between the actuals and predicted values.

$$MSE = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$

Where, \hat{Y}_i is a vector of n predictions and Y is the vector of observed values of the variable being predicted.

Mean Absolute Error(MAE)

The mean absolute error(MAE) measures the mean of the absolute errors that is, the absolute value of the difference between the forecasted value and the actual value. MAE tells us how big of an

error we can expect from the forecast on average.

$$MAE = \frac{1}{n} \sum_{i=1}^n |Y_i - \hat{Y}_i|$$

Where, \hat{Y}_i is a vector of n forecasts and Y is the vector of actual values of the variable being predicted.

MSE Vs MAE

Mean squared error has the disadvantage of heavily weighting outliers. It is a result of the squaring of each term, which effectively weights large errors more heavily than small ones. Where this kind of property is undesirable, MAE can be used in those applications by the researcher.

When dealing with outliers, it might be helpful to use MAE instead of MSE since MSE gives higher error than MAE. Yet, MSE is more popular and efficient than MAE, because MSE punishes larger errors, which tends to be useful in the real world.

The mean absolute error (MAE) has the same unit as the original data, and it can only be compared between models whose errors are measured in the same units.

Both MSE and MAE are scale-dependent. For instance, if the observed data are in km then MSE is in km^2 and MAE is always in km respectively. Often, we need to perform accuracy test on predicted values across different units. In that particular context, both MSE and MAE will not be applicable because they can only be compared between models whose errors are measured in the same units.

For evenly distributed errors that is, when all of the errors have the same magnitude, then Root mean squared error (RMSE) and Mean absolute error (MAE) will give the same result. If the square of the difference between actual values and forecasted values gives a positive distance which is same as their absolute distance then, $MSE = MAE$.

Data collection and exploration

To calculate MSE and MAE of different regression methods we used the *Energy_efficiency.csv* dataset. This dataset has been collected from the UCI Machine Learning Repository^[3]. This dataset is a collection of 768 samples and 8 features, aiming to predict two real valued responses.

The dataset contains the following eight attributes or features (X_1, X_2, \dots, X_8) along with two response variables (Y_1, Y_2):

- Relative Compactness (X_1)
- Surface Area (X_2)
- Wall Area (X_3)
- Roof Area (X_4)
- Overall Height (X_5)
- Orientation (X_6)
- Glazing Area (X_7)
- Glazing Area Distribution (X_8)
- Heating Load (Y_1)
- Cooling Load (Y_2)

It is important to implement energy efficiency in building to mitigate the impact of climate change. Due to the high demand for energy and unsustainable supplies, energy efficiency in building plays a vital role reducing energy costs and greenhouse gas emissions. Therefore, studying this dataset to evaluate how well energy is being used there to cut out the costs which will be helpful to have a ECO-friendly environment.

Experiment and perform evaluation

We load the samples into a dataframe and took all the column attributes as factor. We randomize the data frame using `.sample()`. Then, we divided the dataset into a trained dataset with the top 80% of the samples, and a tested dataset with the bottom 20% of the samples respectively. So, energy train data has first 614 entries from the dataset and energy test data contains the rest 154 samples.

At first we set up a model(*rt1*) for tree regression using the *Heating.Load* as outcome variable and all the eight attributes as input variables and fit a new dataframe with the actual and predicted value of the model based on the test data. Using Regression Tree model(*rt1*) and “Heating.Load” as outcome, we calculated $MSE = 6.59$ and $MAE = 2.101$.

Similarly we fit another model(*rt2*) for tree regression but instead of using *Heating.Load* as outcome variable now we are interested to use *Cooling.Load* as outcome variable. And we figured out for this model(*rt2*), using *Cooling.Load* we got $MSE = 8.461$ and $MAE = 2.084$.

We randomize the data frame using again. Next up we fit two models namely *rf1* and *rf2* respectively for both *Heating.Load* and *Cooling.Load* as outcome variables using Random forest regression following the same approach as described earlier for *rt1* and *rt2*. Then we measured the MSE and MAE and for *rf1* we got, $MSE = 1.36$ and $MAE = 0.907$.

##	% Inc MSE
## Glazing.Area	74.7
## Glazing.Area.Distribution	41.0
## Relative.Compactness	24.7
## Surface.Area	24.1
## Wall.Area	20.8
## Roof.Area	18.9
## Overall.Height	17.7
## Orientation	-19.6

Observing the result of *importance()* function to calculate the importance of each variable, we got to see that *Glazing.Area* was considered the most important predictor; it is estimated that, in the absence of that variable, the error would increase by 74.7%.

Whereas for model *rf2*, using *Cooling.Load* we got $MSE = 3.698$ and $MAE = 1.348$.

##	% Inc MSE
## Glazing.Area	75.51
## Glazing.Area.Distribution	29.36
## Relative.Compactness	24.03
## Surface.Area	22.02
## Roof.Area	21.45
## Wall.Area	20.14

```
## Overall.Height      18.78
## Orientation         -4.36
```

If we look into the *importance()* function to calculate the importance of each variable, we can see that The *Glazing.Area* was considered the most important predictor for *rf2*. it is estimated that, in the absence of that variable, the error would increase by 75.51%.

If we perform overall evaluation and compare MSE and MAE for all four models we can see that using random forest regression the model, *rf1* with *Heating.Load* as response variable has lower error rate for both $MSE = 1.36$ and $MAE = 0.907$ compared to other models. For regression tree model both *rt1* and *rt2* produced relatively higher MSE values though MAE values did not vary significantly.

Task 3.C – Theory: Analyze a less obvious dataset

Theory: Analyze a less obvious dataset : SMS Spam Filtering

Text message classification requires supervised Natural language processing techniques to filter messages with respect to its types and maps inputs to its targeted variables based on the learning information which it gets from trained data.

Our aim is to predict the probabilities of a message being spam or ham. Therefore, we need to perform text mining on unstructured data, fit a predictive model on top of that and suggest improvement if any to increase our proposed model's performance.

Data Collection


The dataset: *SmsCollection.csv* has been collected from the course website. This dataset is a collection of 5574 text messages in English provided by a UK forum for research purpose. In this dataset, messages are labeled as either *spam* or *ham*. *Ham* stands for legitimate message whereas the type *spam* is used for trashed or unwanted message.

At first we load the data from the source. Then we split label and text and bind them into a dataframe.

Data exploration

The *SmsCollection* dataset contains text messages only. Since we are only dealing with text messages which are unstructured in nature, so we will need to perform some basic natural language processing technique in order to tokenize those texts, computing the frequencies of words, calculating document-feature matrix and so on.

In general, almost all the classifiers use a conditional probability model to classify data. Looking at the samples we can see that they are mainly concerning about classifying the messages into a two class problem as spam or ham. Among 5574 text messages there are 4827 messages categorized as ham and the rest 747 messages are classified as spam. We generate a barplot of it.



As we can observe there are more ham messages than spam. There are various classifier algorithms to solve this but we found Naive Bayes as the most suitable one for this purpose. Naive Bayes is a simple yet powerful classifier based on Bayes probability theorem which uses conditional probability model. It is more suited to categorical variables although it can also be used for continuous variables. Text messages are often noisy and the amount of predictors are way more than the actual samples. Naive Bayes classifier follows conditional independence theorem. Therefore, it assumes that features are independent of one another which is a high bias and this introduced strong bias might be helpful in reducing the variance to achieve better predictions.

Data Processing and transformation

We load the samples into a dataframe and use the *label* as a factor while on the otherhand we are using attribute *text* as character. And then we randomize the data frame using `sample_n()`. To process the text data we transformed the data frame into a volatile corpus as they cannot be directly handled by a data frame. VCorpus converted each of the messages as a document.

In the VCorpus text document each SMS has its content in raw formatted way. So, before applying Naive Bayes classification algorithm we need to clean up data. It will help the algorithm to perform more efficiently which will eventually increase the accuracy of predictions.

Our data cleaning process includes : conversion of all texts to lowercase, removal of numbers that is neither a spam nor a ham, removal of some common stop words in english such as: “a”, “an”, “the”, “for” etc. that neither indicate spam or ham, punctuation and extra whitespace removal. Finally after completing data cleaning task, final version of the VCorpus were transformed into a Document-Term-Matrix (DTM) that will be taken into account as the basis for the classification. Doing so, we found 7713 unique terms in total for all 5574 entries.

Generating training and testing dataset

We divided the DTM to generate our training and testing dataset. The Document-term-matrix is splitted into a trained dataset with the top 75% of the raw sms data, and a tested dataset with the bottom 25% of the raw sms data using the *createDataPartition()* function. Since, we only need “label” attribute of the raw sms dataset we created two classifier labels namely “sms train labels” and “sms test labels” by splitting with exact same proportions of row that we used before. We made these two classifier labels to use them for Naive Bayes model later on.

Table 1: Frequency comparison among different datasets based on SMS label

	Raw Dataset	Training Dataset	Test Dataset
ham	86.6	86.6	86.6
spam	13.4	13.4	13.4

Using `prop.table()` we converted number of spam/ham messages of both sms train and test labels into fractional values and preserved those proportions into our train and test dataset. Looking into the above table we can see that 86.6% of the messages correspond to legitimate messages (ham) and 13.4% to spam messages which follows the same proportion in each of our dataset perfectly.

We created a wordcloud from the cleaned vcorpus to look at the most frequent words in the available bag of words. We also created separate wordclouds for spam and ham messages where most frequent words appeared in larger font and less frequent words in smaller font.



Looking the above wordclouds we found that the spam contains “call”, “now”, “free”, “mobile” as most frequent words whereas ham contains frequent words such as “will”, “get”, “now”, “just”, “can”. Also spam(on left) showed extreme frequency in it’s wordcloud. Since, it seemed that our datasets contains distintive words, we hope our choosen classifier algorithm(Naive Bayes) will be a good fit

for sms prediction.

We removed most of the least frequent words from the DTM and created final train and test dataset that we should be using for the training using only the most frequent words that appeared at least 5 times in datasets. The number of columns for each trained and tested datasets are then shrink from 7713 terms to 1193 column (words).

Training the data using Naive Bayes Model

we already have trained and tested labels respective to the datasets. And we used `naive_bayes()` to train and build model based on the trained dataset along with it's trained label. Our trained model contained information from both trained and tested DTM which have 1193 distinct words (possibilities of either spam/ham).

Evaluate performance

evaluating it's performance we can see that Naive Bayes has accuracy rate 96.77% with sensitivity 78.49% and there are 5 *spam* text messages wrongly classified as ham and 40 *ham* text examples wrongly classified as spam. In order to improve it's performance we used Laplace along with it and laplace lowered both of the false positive and false negative values and increased our accuracy up to 97.56%. The summarised version of their performances are given below into tabular form.

Table 2: Performance Table for two models

	True_Neg	True_Pos	False_Neg	False_Pos	accuracy	sensitivity
Model-1:[Naive Bayes]	1201	151	35	5	0.9713	0.8118
Model-2:[Naive Bayes+laplace]	1204	163	23	2	0.982	0.8763

Conclusion

To solve this task we classified text messages as ham or spam using some basic natural language processing and then model a naive Bayes text classifier. There are numerous ways of doing this but using the Naive Bayes classification algorithm, we obtained more than 97% accuracy in predicting whether a new incoming message is a spam or not based on it's training data.

References

1. https://en.wikipedia.org/wiki/Naive_Bayes_classifier
2. https://en.wikipedia.org/wiki/Naive_Bayes_spam_filtering
3. <https://www.r-bloggers.com/understanding-naive-bayes-classifier-using-r/>
4. <https://datascienceplus.com/text-message-classification/>
5. https://en.wikipedia.org/wiki/Mean_squared_error
6. https://en.wikipedia.org/wiki/Mean_absolute_error
7. <http://archive.ics.uci.edu/ml/datasets/Energy+efficiency#>

8. A. Tsanas, A. Xifara: ‘Accurate quantitative estimation of energy performance of residential buildings using statistical machine learning tools’, *Energy and Buildings*, Vol. 49, pp. 560-567, 2012 (the paper can be accessed from weblink)